How to deploy the SPIRAL website on an Azure ubuntu 22.04 server

July 26, 2023

1 How to move SPIRAL to a new Azure Ubuntu 22.04 server?

A helpful tutorial: https://www.rosehosting.com/blog/how-to-install-flask-on-ubuntu-22-04-with-apache-and-wsgi/.

Note that we store the app at /var/www/html/.

1.1 Create a VM and open the 80 (http) and 443 (https) ports (and 5000)

This tutorial may be helpful for this step: https://www.youtube.com/watch?v=HmKUCPRKNos&ab_channel=V. In the Azure portal:

- Follow the tutorial to create a VM. The current deployment uses a Standard E20as v5 (20 vcpus, 160 GiB memory).
- \bullet In the Azure portal: VM -> Networking -> Add in bound security rule -> port 80, name it "allow HTTP"
- \bullet In the Azure portal: VM -> Networking -> Add in bound security rule -> port 443, name it "allow HTTPS"
- In the Azure portal: VM -> Networking -> Add inbound security rule -> port 5000, name it "allow5000" (only for now)
- In the Azure portal: VM -> Networking -> Make sure that SSH is open for everywhere (only for now)

Important: before beginning the installation, make sure that all the disks are allocated and mounted. The current deployment uses a 1T disk that is used for both the OS and data. However, the more "correct" way is to use a small (32G) OS disk, and a larger (let's say 1T) Data disk. In this case, you should mount the Data disk on some folder, let's say /var2, and then create a symbolic link from /var to /var2.

Some useful commands:

- fdisk -1 shows both the physical and attached disks.
- 1sblk shows the physical disks and whether they are partitioned and mounted. So if there is "sdb" for example but it is not divided to "sdb1", "sdb14" etc., then it is not partitioned. If it is divided, but for example "sdb1" has nothing under "MOUNTPOINTS", then it is not mounted.

• df -h shows the actual attached disks (and the available space). So make sure that the "SIZE" fits what you expect.

1.2 Apache2

Open Putty and connect to the VM (public IP) with you name and password.

As in the tutorial:

Before installing the software, we need to update the system packages to the latest versions available.

```
sudo apt-get update -y && sudo apt-get upgrade -y
```

To install the Apache Web server execute the following command:

```
sudo apt install apache2 -y
```

Once installed, start and enable the service.

sudo systemctl enable apache2 && sudo systemctl start apache2

Check if the service is up and running:

```
sudo systemctl status apache2
```

Open the server IP in a browser. You should see the apache2 ubuntu default page.

1.3 Python

Verify Python installation:

```
python3 -V
```

you should see Python 3.10.x

Install python3-pip, python3-venv, python3-dev:

sudo apt-get install python3-pip python3-venv python3-dev

1.4 Clone SPIRAL repository, install packages and add needed files

cd /var/www/html/

As explained in the readme of https://github.com/hadasbi/SPIRAL.web.tool (steps 2-7):

```
sudo git clone https://github.com/hadasbi/SPIRAL.web.tool.git
cd SPIRAL.web.tool
sudo python3 -m venv spiral venv
```

Activate the virtual environment:

source ./spiral_venv/bin/activate

Install required Python packages:

sudo ./spiral_venv/bin/pip install -r ./requirements.txt

Create analysis directory:

cd static
sudo mkdir analysis
cd ..

Move required files (that are not in Github for obvious reasons...) to the server: The files are Flask-WTF_encription_key.txt and mail_password.txt.

In winSCP, move the files from your computer to /home/SpiralAdmin (since you can only write to /var/www/html in sudo). Then move them to their final location by:

```
sudo mv /home/SpiralAdmin/mail_password.txt ./mail_password.txt
sudo mv /home/SpiralAdmin/Flask-WTF_encription_key.txt ./Flask-WTF_encription_key.txt
```

Add the server name to list of servers running in production mode (in main.py) In main.py, look for the line:

```
if hostname in ['bi-fiona', 'Spin']: # production (linux)
```

And add <SERVERNAME>:

```
if hostname in ['bi-fiona', 'Spin', '<SERVERNAME>']: # production (linux)
```

Check that the flask app is working: A very important note: app.config['SERVER_NAME'] in \var\www\html\SPIRAL.web.tool\main.py and ServerName in the configuration file /etc/apache2/sites-available/spiral.technion.ac.il.conf (see next section) have to match.

For now, in order to test the app, change app.config['SERVER_NAME'] in \var\www\html\SPIRAL.web.tool\main.py to <VM_PUBLIC_IP>:5000 with the command:

sudo nano \var\www\html\SPIRAL.web.tool\main.py

Navigate the file and make the change.

To save, type: ctrl+X, then Y, then enter.

Then you can test the app:

Test 1: make sure you have the venv activated and type:

python main.py

Test 2: make sure you have the venv activated and type:

```
export FLASK_APP=main.py
flask run --host=0.0.0.0
```

For both tests you should see:

Running on all addresses.

WARNING: This is a development server. Do not use it in a production deployment.

Running on http://10.0.0.8:5000/ (Press CTRL+C to quit)

Open the browser and type: <VM_PUBLIC_IP>:5000

You should see SPIRAL's homepage.

Go back to the terminal, press ctrl+C and:

deactivate

1.5 What if it doesn't work?

Go to the last line in main.py and add debug=True to app.run().

Then:

python file.py runserver -d

2 WSGI configuration for HTTP

This step is only important for Windows, not linux:

sudo cp activate_this.py ./spiral_venv/bin/activate_this.py

Install mod_wsgi:

sudo apt-get install libapache2-mod-wsgi-py3

Create SPIRAL configuration file:

sudo touch /etc/apache2/sites-available/spiral.technion.ac.il.conf

To edit the file, open it with nano:

sudo nano /etc/apache2/sites-available/spiral.technion.ac.il.conf

And paste this:

<VirtualHost *:80>

#ServerName spiral.technion.ac.il

ServerName <VM_PUBLIC_IP>

ServerAlias www.spiral.technion.ac.il

DocumentRoot "/var/www/html/SPIRAL.web.tool/"

WSGIDaemonProcess spiral user=www-data group=www-data threads=20 python-home="/var/www/html/SP #python-path=/var/www/html/SPIRAL.web.tool/:/var/www/html/SPIRAL.web.tool/spiral_venv/lib/pythowsGIScriptAlias / "/var/www/html/SPIRAL.web.tool/app.py"

ErrorLog /var/log/apache2/spiral.technion.ac.il-error.log
CustomLog /var/log/apache2/spiral.technion.ac.il-access.log combined

<Directory "/var/www/html/SPIRAL.web.tool">

WSGIProcessGroup spiral
WSGIApplicationGroup %{GLOBAL}
Order deny,allow
Require all granted
</Directory>
LogLevel warn
</VirtualHost>

To save, type: ctrl+X, then Y, then enter.

Then, enable the Apache2 configuration file:

sudo a2ensite spiral.technion.ac.il.conf

Check the syntax of the Apache2 configuration:

apachectl -t

You should see:

Syntax OK

Restart Apache:

sudo systemctl restart apache2

Remove index.html (the default Apache page) from /var/www/html/:

sudo mv /var/www/html/index.html /home/SpiralAdmin/index.html

Change the ownership of the relevant folder to the user that is supposed to run the app:

sudo chown -R www-data:www-data /var/www/

Change server name: Change both app.config['SERVER_NAME'] in \var\www\html\SPIRAL.web.tool\main.py and ServerName in the configuration file /etc/apache2/sites-available/spiral.technion.ac.il.conf to <VM_PUBLIC_IP>.

Check if it's working: Restart Apache:

sudo systemctl restart apache2

Open http://<VM_PUBLIC_IP>/ in a browser. SPIRAL's homepage should appear.

3 WSGI configuration for HTTPS

Follow this tutorial: https://www.digitalocean.com/community/tutorials/how-to-create-a-self-signed-ssl-certificate-for-apache-in-ubuntu-22-04,

but change the conf file /etc/apache2/sites-available/spiral.technion.ac.il.conf to:

<VirtualHost *:80>
ServerName <VM_PUBLIC_IP>
ServerAlias www.spiral.technion.ac.il

Redirect / https://<VM_PUBLIC_IP>/
</VirtualHost>

<VirtualHost *:443>
#ServerName spiral.technion.ac.il
ServerName <VM_PUBLIC_IP>
ServerAlias www.spiral.technion.ac.il
DocumentRoot "/var/www/html/SPIRAL.web.tool/"

WSGIDaemonProcess spiral user=www-data group=www-data processes=19 maximum-requests=100 python-python-path=/var/www/html/SPIRAL.web.tool/:/var/www/html/SPIRAL.web.tool/spiral_venv/lib/python-

ErrorLog /var/log/apache2/spiral.technion.ac.il-error.log CustomLog /var/log/apache2/spiral.technion.ac.il-access.log combined

<Directory "/var/www/html/SPIRAL.web.tool">

WSGIProcessGroup spiral
WSGIApplicationGroup %{GLOBAL}
Order deny,allow
Require all granted
</Directory>
LogLevel warn

SSLEngine on

#SSLCertificateFile /etc/ssl/certs/wildcard.technion.ac.il.crt
#SSLCertificateKeyFile /etc/ssl/private/wildcard.technion.ac.il.key
#SSLCACertificateFile /etc/ssl/certs/technion-ca-bundle.crt
SSLCertificateFile /etc/ssl/certs/apache-selfsigned.crt
SSLCertificateKeyFile /etc/ssl/private/apache-selfsigned.key
</VirtualHost>

- Ask Dora for the Technion certificate key for the computer (she needs <VM_PUBLIC_IP>). Note that the server time response (the time that it takes until homepage appears) may be long while it still relies on the self assigned certificate.
- Ask that the spiral.technion.ac.il DNS would be moved to <VM_PUBLIC_IP>.

4 When you have the Technion certificates:

Copy them to the right place (see conf file) and change the conf file accordingly.

5 After DNS has been moved:

Change server name: Change both app.config['SERVER_NAME'] in \var\www\html\SPIRAL.web.tool\main.py and ServerName in the configuration file /etc/apache2/sites-available/spiral.technion.ac.il.conf to spiral.technion.ac.il.

6 Some last tasks...

- In the Azure portal: remove the port 5000 rule.
- In the Azure portal: Make sure that SSH is only open for a list of IPs (enter your IP). Remember this... next time you want to connect via SSH you'll need to first go to the Azure portal and enter your IP in the list of allowed IPs.

7 Common issues...

7.0.1 What to do if something isn't working?

```
Check the error logs (both SPIRAL's and Apache's):

cat /var/log/apache2/spiral.technion.ac.il-error.log

cat /var/log/apache2/error.log
```

Azure portal: make sure that your NIC and your subnet are under the same NSG.

Don't install nginx, as it can create issues by making the ports unavaliable for Apache2. If nginx is installed you have to completely remove it (see https://www.cyberciti.biz/faq/remove-uninstall-nginx-from-ubuntu-debian-linux-apt-get-command/).

```
[]:
```

```
[13]: # Check response time
import requests, time

url = 'https://4.227.224.198/'
url = 'http://4.227.224.198/'
#url = 'http://20.127.4.198/'

t0 = time.time()
response = requests.get(url, verify=False)
t1 = time.time()
total = t1-t0
print("Simple get request took " , total)
```

c:\program files (x86)\microsoft visual studio\shared\python37_64\lib\site-packages\urllib3\connectionpool.py:1020: InsecureRequestWarning: Unverified HTTPS request is being made to host '4.227.224.198'. Adding certificate verification is strongly advised. See:

https://urllib3.readthedocs.io/en/1.26.x/advanced-usage.html#ssl-warnings InsecureRequestWarning,

Simple get request took 10.989382266998291